How To Install Nginx on Ubuntu 18.04

Posted **April 27, 2018** 👁 171.9k   NGINX   UBUNTU 18.04

By: Justin Ellingwood   By: Kathleen Juell

Not using **Ubuntu 18.04**? Choose a different version:

Introduction

Nginx is one of the most popular web servers in the world and is responsible for hosting some of the largest and highest-traffic sites on the internet. It is more resource-friendly than Apache in most cases and can be used as a web server or reverse proxy.

In this guide, we'll discuss how to install Nginx on your Ubuntu 18.04 server.

Prerequisites

SCROLL TO TOP

Before you begin this guide, you should have a regular, non-root user with sudo privileges configured on your server. You can learn how to configure a regular user account by following our initial server setup guide for Ubuntu 18.04.

When you have an account available, log in as your non-root user to begin.

Step 1 – Installing Nginx

Because Nginx is available in Ubuntu's default repositories, it is possible to install it from these repositories using the `apt` packaging system.

Since this is our first interaction with the `apt` packaging system in this session, we will update our local package index so that we have access to the most recent package listings. Afterwards, we can install `nginx`:

```
$ sudo apt update
$ sudo apt install nginx
```

After accepting the procedure, `apt` will install Nginx and any required dependencies to your server.

Step 2 – Adjusting the Firewall

Before testing Nginx, the firewall software needs to be adjusted to allow access to the service. Nginx registers itself as a service with `ufw` upon installation, making it straightforward to allow Nginx access.

List the application configurations that `ufw` knows how to work with by typing:

```
$ sudo ufw app list
```

You should get a listing of the application profiles:

```
Output
Available applications:
  Nginx Full
  Nginx HTTP
  Nginx HTTPS
  OpenSSH
```

As you can see, there are three profiles available for Nginx:

- **Nginx Full**: This profile opens both port 80 (normal, unencrypted web traffic) and port 443 (TLS/SSL encrypted traffic)

- **Nginx HTTP**: This profile opens only port 80 (normal, unencrypted web traffic)

- **Nginx HTTPS**: This profile opens only port 443 (TLS/SSL encrypted traffic)

It is recommended that you enable the most restrictive profile that will still allow the traffic you've configured. Since we haven't configured SSL for our server yet in this guide, we will only need to allow traffic on port 80.

You can enable this by typing:

```
$ sudo ufw allow 'Nginx HTTP'
```

You can verify the change by typing:

```
$ sudo ufw status
```

You should see HTTP traffic allowed in the displayed output:

Output
```
Status: active

To                         Action      From
--                         ------      ----
OpenSSH                    ALLOW       Anywhere
Nginx HTTP                 ALLOW       Anywhere
OpenSSH (v6)               ALLOW       Anywhere (v6)
Nginx HTTP (v6)            ALLOW       Anywhere (v6)
```

Step 3 – Checking your Web Server

At the end of the installation process, Ubuntu 18.04 starts Nginx. The web server should already be up and running.

We can check with the `systemd` init system to make sure the service is running by typing:

```
$ systemctl status nginx
```

Output
```
● nginx.service - A high performance web server and a reverse proxy server
   Loaded: loaded (/lib/systemd/system/nginx.service; enabled; vendor preset: enabled)
   Active: active (running) since Fri 2018-04-20 16:08:19 UTC; 3 days ago
     Docs: man:nginx(8)
 Main PID: 2369 (nginx)
    Tasks: 2 (limit: 1153)
   CGroup: /system.slice/nginx.service
           └─2369 nginx: master process /usr/sbin/nginx -g daemon on; master process on;
```

As you can see above, the service appears to have started successfully. However, the best way to test this is to actually request a page from Nginx.

You can access the default Nginx landing page to confirm that the software is running properly by navigating to your server's IP address. If you do not know your server's IP address, you can get it a few different ways.

Try typing this at your server's command prompt:

```
$ ip addr show eth0 | grep inet | awk '{ print $2; }' | sed 's/\/.*$//'
```

You will get back a few lines. You can try each in your web browser to see if they work.

An alternative is typing this, which should give you your public IP address as seen from another location on the internet:

```
$ curl -4 icanhazip.com
```

When you have your server's IP address, enter it into your browser's address bar:

```
http://your_server_ip
```

You should see the default Nginx landing page:

This page is included with Nginx to show you that the server is running correctly.

Step 4 – Managing the Nginx Process

Now that you have your web server up and running, let's review some basic management commands.

To stop your web server, type:

```
$ sudo systemctl stop nginx
```

To start the web server when it is stopped, type:

```
$ sudo systemctl start nginx
```

To stop and then start the service again, type:

```
$ sudo systemctl restart nginx
```

If you are simply making configuration changes, Nginx can often reload without dropping connections. To do this, type:

```
$ sudo systemctl reload nginx
```

By default, Nginx is configured to start automatically when the server boots. If this is not what you want, you can disable this behavior by typing:

```
$ sudo systemctl disable nginx
```

To re-enable the service to start up at boot, you can type:

```
$ sudo systemctl enable nginx
```

Step 5 – Setting Up Server Blocks (Recommended)

When using the Nginx web server, *server blocks* (similar to virtual hosts in Apache) can be used to encapsulate configuration details and host more than one domain from a single server. We will set up a domain called **example.com**, but you should **replace this with your own domain name**. To learn more about setting up a domain name with DigitalOcean, see our Introduction to DigitalOcean DNS.

Nginx on Ubuntu 18.04 has one server block enabled by default that is configured to serve documents out of a directory at `/var/www/html`. While this works well for a single site, it can become unwieldy if you are hosting multiple sites. Instead of modifying `/var/www/html`, let's create a directory structure within `/var/www` for our **example.com** site, leaving `/var/www/html` in place as the default directory to be served if a client request doesn't match any other sites.

Create the directory for **example.com** as follows, using the `-p` flag to create any necessary parent directories:

```
$ sudo mkdir -p /var/www/example.com/html
```

Next, assign ownership of the directory with the `$USER` environment variable:

```
$ sudo chown -R $USER:$USER /var/www/example.com/html
```

The permissions of your web roots should be correct if you haven't modified your `umask` value, but you can make sure by typing:

```
$ sudo chmod -R 755 /var/www/example.com
```

Next, create a sample `index.html` page using `nano` or your favorite editor:

```
$ nano /var/www/example.com/html/index.html
```

Inside, add the following sample HTML:

/var/www/example.com/html/index.html

```html
<html>
    <head>
        <title>Welcome to Example.com!</title>
    </head>
    <body>
        <h1>Success!  The example.com server block is working!</h1>
    </body>
</html>
```

Save and close the file when you are finished.

In order for Nginx to serve this content, it's necessary to create a server block with the correct directives. Instead of modifying the default configuration file directly, let's make a new one at `/etc/nginx/sites-available/example.com`:

```
$ sudo nano /etc/nginx/sites-available/example.com
```

Paste in the following configuration block, which is similar to the default, but updated for our new directory and domain name:

/etc/nginx/sites-available/example.com

```
server {
        listen 80;
        listen [::]:80;

        root /var/www/example.com/html;
```

```
        index index.html index.htm index.nginx-debian.html;


        server_name example.com www.example.com;


        location / {
                try_files $uri $uri/ =404;
        }
}
```

Notice that we've updated the `root` configuration to our new directory, and the `server_name` to our domain name.

Next, let's enable the file by creating a link from it to the `sites-enabled` directory, which Nginx reads from during startup:

```
$ sudo ln -s /etc/nginx/sites-available/example.com /etc/nginx/sites-enabled/
```

Two server blocks are now enabled and configured to respond to requests based on their `listen` and `server_name` directives (you can read more about how Nginx processes these directives here):

- `example.com` : Will respond to requests for `example.com` and `www.example.com` .

- `default` : Will respond to any requests on port 80 that do not match the other two blocks.

To avoid a possible hash bucket memory problem that can arise from adding additional server names, it is necessary to adjust a single value in the `/etc/nginx/nginx.conf` file. Open the file:

```
$ sudo nano /etc/nginx/nginx.conf
```

Find the `server_names_hash_bucket_size` directive and remove the `#` symbol to uncomment the line:

/etc/nginx/nginx.conf

```
...
http {
    ...
    server_names_hash_bucket_size 64;
    ...
}
...
```

Next, test to make sure that there are no syntax errors in any of your Nginx files:

```
$ sudo nginx -t
```

Save and close the file when you are finished.

If there aren't any problems, restart Nginx to enable your changes:

```
$ sudo systemctl restart nginx
```

Nginx should now be serving your domain name. You can test this by navigating to `http://example.com`, where you should see something like this:

Success! The example.com server block is working!

Step 6 – Getting Familiar with Important Nginx Files and Directories

Now that you know how to manage the Nginx service itself, you should take a few minutes to familiarize yourself with a few important directories and files.

Content

- `/var/www/html` : The actual web content, which by default only consists of the default Nginx page you saw earlier, is served out of the `/var/www/html` directory. This can be changed by altering Nginx configuration files.

Server Configuration

- `/etc/nginx` : The Nginx configuration directory. All of the Nginx configuration files reside here.

- `/etc/nginx/nginx.conf` : The main Nginx configuration file. This can be modified to make changes to the Nginx global configuration.

- `/etc/nginx/sites-available/` : The directory where per-site server blocks can be stored. Nginx will not use the configuration files found in this directory unless they are linked to the `sites-enabled` directory. Typically, all server block configuration is done in this directory, and then enabled by linking to the other directory.

- `/etc/nginx/sites-enabled/` : The directory where enabled per-site server blocks are stored. Typically, these are created by linking to configuration files found in the `sites-available` directory.

- `/etc/nginx/snippets` : This directory contains configuration fragments that can be included elsewhere in the Nginx configuration. Potentially repeatable configuration segments are good candidates for refactoring into snippets.

Server Logs

- `/var/log/nginx/access.log` : Every request to your web server is recorded in this log configured to do otherwise.

- `/var/log/nginx/error.log` : Any Nginx errors will be recorded in this log.

Conclusion

Now that you have your web server installed, you have many options for the type of content to serve and the technologies you want to use to create a richer experience.

If you'd like to build out a more complete application stack, check out this article on how to configure a LEMP stack on Ubuntu 18.04.

By: Justin Ellingwood   By: Kathleen Juell   ♡ Upvote (53)   ⊡ Subscribe   ⬆ Share

We just made it easier for you to deploy faster.

TRY FREE

Related Tutorials

How To Upgrade Nginx In-Place Without Dropping Client Connections

How To Target Your Users with Nginx Analytics and A/B Testing

How To Deploy a PHP Application with Kubernetes on Ubuntu 16.04

How To Ensure Code Quality with SonarQube on Ubuntu 18.04

How To Set Up a Private Docker Registry on Ubuntu 18.04

10 Comments

Leave a comment...

Log In to Comment

rgsives *July 2, 2018*

0 Will I be able to access Radio Romania International? Their address is www.rri.ro

strobingRaptor *July 11, 2018*

3 This works flawlessly. Very clear guide. Thankyou!

nathan4cs *July 16, 2018*

1 After step 2, I got `Status: inactive` and discovered that I needed to run `sudo ufw enable`. And after that the output was a bit different than described:

```
Status: active

To                        Action      From
--                        ------      ----
Nginx HTTP                ALLOW       Anywhere
Nginx HTTP (v6)           ALLOW       Anywhere (v6)
```

Lami *July 18, 2018*

1 I've had same problem, you just need to type `sudo ufw allow 'OpenSSH` and that's all

mokoshiba *August 29, 2018*

0 Had the same problem, should have been mentioned in the article

neoNitish *September 18, 2018*

0 I am also getting the same problem. Did it work after `sudo ufw enable` command?

SCROLL TO TOP

neoNitish *September 18, 2018*

0 I was also getting `Status: inactive` by typing `sudo ufw allow 'Nginx HTTP'`. But I solved it by using `sudo ufw enable` command.

niyazitoros *October 8, 2018*

0 Hi,

On vmware I have Ubuntu 18.04. When I try this "ip addr show eth0 | grep inet | awk '{ print $2; }' | sed 's/\/.*$//'" I got error saying "Device "eth0" does not exist." My DNS is working. I also follow the tutorial and create block for domin. But I cannot make it work. when I type http://127.0.0.1/ I can see default nginx index.html But when I try to write my domain it looks like its keep searching. Any idea?

niyazi@niyazi-virtual-machine:~$ sudo lshw -short | grep network
/0/100/15/0 ens160 network VMXNET3 Ethernet Controller
/2 docker0 network Ethernet interface

cyntss *November 20, 2018*

0 Very good guidelines.

niyazitoros *December 4, 2018*

0 I follow the tutorial and set up my Nginx as well as I configure to use as a bypass proxy. When I try to type as http://myserverip/ I get 502 bad gateway error. I am using Nginx bypass proxy so I can test my web api. I created single customize html and put in root /var/www/html/data/; but I still keep getting 502 bad gateway error. How can I correct my configuration?

Below is my Nginx configuration. Any help please?

```
    server {
            charset UTF-8;
            listen 80;
            listen [::]:80;
            error_log    /var/log/nginx/niyazi.com.error.log debug;
            rewrite_log on;
            server_name mod.niyazi.com;
            include snippets/letsencrypt.conf;
            root /var/www/html/data/;
            return 301 https://mod.niyazi.com$request_uri;

        }

        server {
            charset UTF-8;
            listen 443 ssl http2;
            server_name mod.niyazi.com;
```

```
        ssl_certificate /etc/letsencrypt/live/mod.niyazi.com/fullchain.pem;
        ssl_certificate_key /etc/letsencrypt/live/mod.niyazi.com/privkey.pem;
        ssl_trusted_certificate /etc/letsencrypt/live/mod.niyazi.com/chain.pem;
        include snippets/ssl.conf;
        include snippets/letsencrypt.conf;
        location / {
                proxy_pass http://192.168.12.120:8280/;
        }
        if ($request_method !~ ^(GET|HEAD|POST)$ )
    {
    return 405;
    }
}
```

SCROLL TO TOP